

Bayesian Theory and Computation

Lecture 19: Energy-based and Score-based Generative Models



Cheng Zhang

School of Mathematical Sciences, Peking University

May 14, 2025

Probability densities $p(x)$ need to satisfy

- ▶ non-negative: $p(x) \geq 0$.
- ▶ sum-to-one: $\sum_x p(x) = 1$ or $\int p(x)dx = 1$ for continuous variables

Coming up with a non-negative function $p_\theta(x)$ is not hard

- ▶ $p_\theta(x) = f_\theta(x)^2$
- ▶ $p_\theta(x) = \exp(f_\theta(x))$
- ▶ $p_\theta(x) = |f_\theta(x)|$

Sum to one is the key. Although many models allow analytical integration (e.g., autoregressive models, normalizing flows), what if the analytical integration is not available?



$$p_\theta(x) = \frac{\exp(f_\theta(x))}{Z(\theta)}, \quad Z(\theta) = \int \exp(f_\theta(x))dx$$

The normalizing constant $Z(\theta)$ is also called the partition function. Why exponential (and not e.g. $f_\theta(x)^2$)?

- ▶ Want to capture very large variations in probability.
log-probability is the natural scale we want to work with.
Otherwise need highly non-smooth f_θ .
- ▶ Exponential families. Many common distributions can be written in this form.
- ▶ These distributions arise under fairly general assumptions in statistical physics (maximum entropy, second law of thermodynamics).
 - ▶ $f_\theta(x)$ is called the energy, hence the name.
 - ▶ Intuitively, configurations x with low energy (high $f_\theta(x)$) are more likely.

$$p_{\theta}(x) = \frac{\exp(f_{\theta}(x))}{Z(\theta)}, \quad Z(\theta) = \int \exp(f_{\theta}(x))dx$$

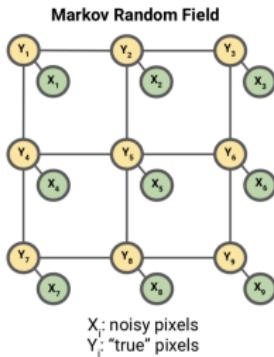
Pros:

- ▶ extreme flexibility. pretty much any function $f_{\theta}(x)$ you want to use

Cons:

- ▶ Sampling from $p_{\theta}(x)$ is hard
- ▶ Evaluating and optimizing likelihood $p_{\theta}(x)$ is hard (learning is hard)
- ▶ No feature learning (but can add latent variables)


Curse of dimensionality: The fundamental issue is that computing $Z(\theta)$ numerically (when no analytic solution is available) scales exponentially in the number of dimensions of x .



- ▶ There is a true image $y \in \{0, 1\}^{3 \times 3}$, and a corrupted image $x \in \{0, 1\}^{3 \times 3}$. We know x , and want to somehow recover y .
- ▶ We model the joint probability distribution $p(y, x)$ as

$$p(y, x) \propto \exp \left(\sum_i \psi_i(x_i, y_i) + \sum_{i,j \in E} \psi_{i,j}(y_i, y_j) \right)$$





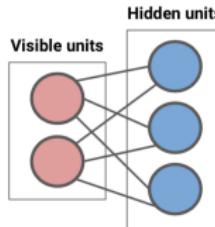
The energy is $\sum_i \psi_i(x_i, y_i) + \sum_{i,j \in E} \psi_{i,j}(y_i, y_j)$

- ▶ $\psi_i(x_i, y_i)$: the i -th corrupted pixel depends on the i -th original pixel
- ▶ $\psi_{i,j}(y_i, y_j)$: neighboring pixels tend to have the same value

How did the original image y look like? Solution: maximize $p(y|x)$. Or equivalently, maximize $p(y, x)$.

- ▶ RBM: energy-based model with latent variables
- ▶ Two types of variables:
 - ▶ $x \in \{0, 1\}^n$ are visible variables (e.g., pixel values)
 - ▶ $z \in \{0, 1\}^m$ are latent ones
- ▶ The joint distribution is

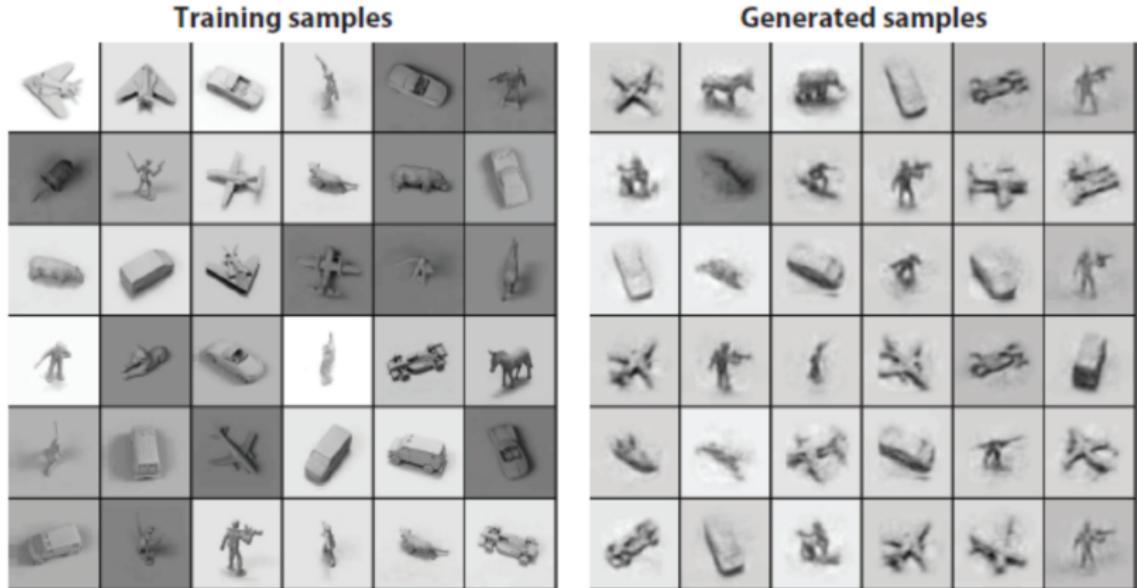
$$p_{W,b,c}(x, z) \propto \exp(x^T W z + b^T x + c^T z)$$



- ▶ Restricted as there are no within-class connections.
- ▶ Can be stacked together to make deep RBMs (one of the first generative models).

Deep RBMs: Samples

7/50



Adapted from Salakhutdinov and Hinton, 2009.

- ▶ Learning by maximizing the likelihood function

$$\max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\theta}(x) = \max_{\theta} (\mathbb{E}_{x \sim p_{\text{data}}} f_{\theta}(x) - \log Z(\theta))$$

- ▶ Gradient of log-likelihood:

$$\begin{aligned}\mathbb{E}_{x \sim p_{\text{data}}} \nabla_{\theta} f_{\theta}(x) - \nabla_{\theta} \log Z(\theta) &= \mathbb{E}_{x \sim p_{\text{data}}} \nabla_{\theta} f_{\theta}(x) - \frac{\nabla_{\theta} Z(\theta)}{Z(\theta)} \\ &= \mathbb{E}_{x \sim p_{\text{data}}} \nabla_{\theta} f_{\theta}(x) - \int \frac{\exp(f_{\theta}(x))}{Z(\theta)} \nabla_{\theta} f_{\theta}(x) dx \\ &= \mathbb{E}_{x \sim p_{\text{data}}} \nabla_{\theta} f_{\theta}(x) - \int p_{\theta}(x) \nabla_{\theta} f_{\theta}(x) dx \\ &= \mathbb{E}_{x \sim p_{\text{data}}} \nabla_{\theta} f_{\theta}(x) - \mathbb{E}_{x \sim p_{\theta}(x)} \nabla_{\theta} f_{\theta}(x)\end{aligned}$$

- ▶ **Contrastive Divergence**: sample $x_{\text{sample}} \sim p_{\theta}$, take gradient step on $\nabla_{\theta} f_{\theta}(x_{\text{train}}) - \nabla_{\theta} f_{\theta}(x_{\text{sample}})$.

$$p_\theta(x) = \frac{\exp(f_\theta(x))}{Z(\theta)}, \quad Z(\theta) = \int \exp(f_\theta(x))dx$$

- ▶ No direct way to sample like in autoregressive or flow models.
- ▶ Can use gradient-based MCMC methods, e.g., SGLD

$$x^{t+1} = x^t + \epsilon \nabla_x \log p_\theta(x^t) + \sqrt{2\epsilon} \eta^t, \quad \eta^t \sim \mathcal{N}(0, I)$$


- ▶ Note that for energy-based models

$$s_\theta(x) = \nabla_x \log p_\theta(x) = \nabla_x f_\theta(x) - \nabla_x \log Z(\theta) = \nabla_x f_\theta(x)$$

The score function does not depend on $Z(\theta)$!



Langevin sampling



Face samples

Adapted from Nijkamp et al. 2019

$$x^{t+1} = x^t + \epsilon \nabla_x \log p_\theta(x^t) + \sqrt{2\epsilon} \eta^t, \quad \eta^t \sim \mathcal{N}(0, I)$$

- ▶ MCMC sampling converges slowly in high dimensional spaces, and repetitive sampling for each training iteration would be expensive.
- ▶ Can we train without sampling?
- ▶ Note that to generate samples from an EBM, we only need the **score function** $\nabla_x \log p_\theta(x)$.
- ▶ Can we properly train the score function without sampling?



- ▶ A key observation: two distributions are identical iff their scores are the same

$$p(x) = q(x) \Leftrightarrow \nabla_x \log \tilde{p}(x) = \nabla_x \log \tilde{q}(x)$$

where \tilde{p}, \tilde{q} are the unnormalized densities of p, q .

- ▶ Match the scores of the data distribution and EBMs by minimizing

$$\begin{aligned} & \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \|\nabla_x \log p_{\text{data}}(x) - s_\theta(x)\|^2 \\ &= \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \|\nabla_x \log p_{\text{data}}(x) - \nabla_x f_\theta(x)\|^2 \end{aligned}$$

This is also known as **Fisher divergence**.



- ▶ Using integration by parts, we have

$$\begin{aligned} & \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \|\nabla_x \log p_{\text{data}}(x) - \nabla_x f_\theta(x)\|^2 \\ &= \mathbb{E}_{x \sim p_{\text{data}}} \left(\text{Tr}(\nabla_x^2 f_\theta(x)) + \frac{1}{2} \|\nabla_x f_\theta(x)\|^2 \right) + \text{Const} \end{aligned}$$

- ▶ Sample a mini-batch of datapoints

$$\{x_1, x_2, \dots, x_n\} \sim p_{\text{data}}(x)$$

- ▶ Estimate the score matching loss with the empirical mean

$$\frac{1}{n} \sum_{i=1}^n \left(\frac{1}{2} \|\nabla_x f_\theta(x_i)\|^2 + \text{Tr}(\nabla_x^2 f_\theta(x_i)) \right)$$

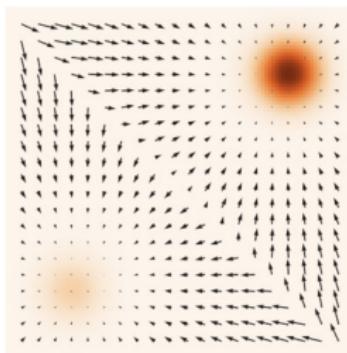


- ▶ Minimize the score matching loss via stochastic gradient descent.
- ▶ No need to sample from the EBM!
- ▶ Note that computing the trace of Hessian $\text{Tr}(\nabla_x^2 f_\theta(x))$ is in general very expensive for large models.
- ▶ Scalable score matching methods: denoising score matching (Vincent 2010) and sliced score matching (Song et al. 2019).

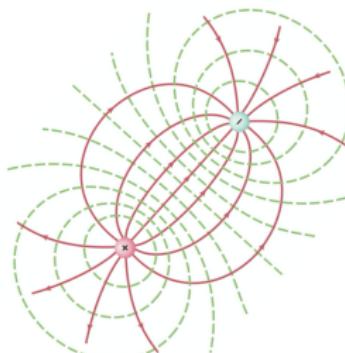


- When the pdf is differentiable, we can compute the gradient of a probability density, and use it to represent the distribution.

Score function $\nabla_x \log p(x)$



(pdf and score)



(Electrical potentials and fields)



- ▶ Given i.i.d. samples $\{x_1, \dots, x_N\} \sim p(x)$
- ▶ We want to estimate the score $\nabla_x \log p_{\text{data}}(x)$
- ▶ Score model: a learnable vector-valued function

$$s_\theta(x) : \mathbb{R}^D \rightarrow \mathbb{R}^D$$

- ▶ Goal: $s_\theta(x) \approx \nabla_x \log p_{\text{data}}(x)$
- ▶ How to compare two vector fields of scores?



- ▶ Objective: Average Euclidean distance over the whole space.

$$\frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \|\nabla_x \log p_{\text{data}}(x) - s_\theta(x)\|^2$$

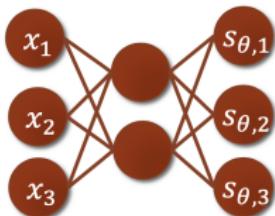
- ▶ **Score matching:**

$$\mathbb{E}_{x \sim p_{\text{data}}} \left(\frac{1}{2} \|s_\theta(x)\|^2 + \text{Tr}(\nabla_x s_\theta(x)) \right)$$

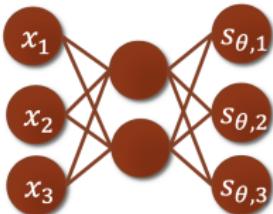
- ▶ Requirements:
 - ▶ The score model must be efficient to evaluated.
 - ▶ Do we need the score model to be a proper score function?



- We can use deep neural networks for more expressive score models



- We can use deep neural networks for more expressive score models



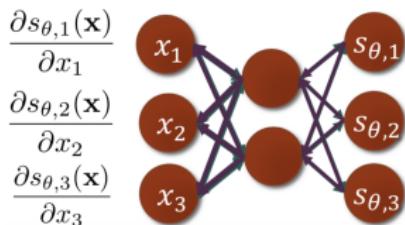
- However, $\text{Tr}(\nabla_x s_\theta(x))$ can be a problem.

$O(D)$ Backprops!

$$\nabla_{\mathbf{x}} s_\theta(\mathbf{x}) = \begin{pmatrix} \frac{\partial s_{\theta,1}(\mathbf{x})}{\partial x_1} & \frac{\partial s_{\theta,1}(\mathbf{x})}{\partial x_2} & \frac{\partial s_{\theta,1}(\mathbf{x})}{\partial x_3} \\ \frac{\partial s_{\theta,2}(\mathbf{x})}{\partial x_1} & \frac{\partial s_{\theta,2}(\mathbf{x})}{\partial x_2} & \frac{\partial s_{\theta,2}(\mathbf{x})}{\partial x_3} \\ \frac{\partial s_{\theta,3}(\mathbf{x})}{\partial x_1} & \frac{\partial s_{\theta,3}(\mathbf{x})}{\partial x_2} & \frac{\partial s_{\theta,3}(\mathbf{x})}{\partial x_3} \end{pmatrix}$$



- We can use deep neural networks for more expressive score models




- However, $\text{Tr}(\nabla_{\mathbf{x}} s_{\theta}(\mathbf{x}))$ can be a problem.

$O(D)$ Backprops!

$$\nabla_{\mathbf{x}} s_{\theta}(\mathbf{x}) = \begin{pmatrix} \frac{\partial s_{\theta,1}(\mathbf{x})}{\partial x_1} & \frac{\partial s_{\theta,1}(\mathbf{x})}{\partial x_2} & \frac{\partial s_{\theta,1}(\mathbf{x})}{\partial x_3} \\ \frac{\partial s_{\theta,2}(\mathbf{x})}{\partial x_1} & \frac{\partial s_{\theta,2}(\mathbf{x})}{\partial x_2} & \frac{\partial s_{\theta,2}(\mathbf{x})}{\partial x_3} \\ \frac{\partial s_{\theta,3}(\mathbf{x})}{\partial x_1} & \frac{\partial s_{\theta,3}(\mathbf{x})}{\partial x_2} & \frac{\partial s_{\theta,3}(\mathbf{x})}{\partial x_3} \end{pmatrix}$$





- ▶ Denoising score matching (Vincent 2011) used a noise-perturbed data distribution

$$\begin{aligned} & \frac{1}{2} \mathbb{E}_{\tilde{x} \sim q_\sigma} \|\nabla_{\tilde{x}} \log q_\sigma(\tilde{x}) - s_\theta(\tilde{x})\|^2 \\ &= \frac{1}{2} \int q_\sigma(\tilde{x}) \|\nabla_{\tilde{x}} \log q_\sigma(\tilde{x}) - s_\theta(\tilde{x})\|^2 d\tilde{x} \\ &= \frac{1}{2} \int q_\sigma(\tilde{x}) \|s_\theta(\tilde{x})\|^2 d\tilde{x} \\ &\quad - \int q_\sigma(\tilde{x}) \nabla_{\tilde{x}} \log q_\sigma(\tilde{x})^T s_\theta(\tilde{x}) d\tilde{x} + \text{Const} \end{aligned}$$

- The second term can be rewritten as

$$\begin{aligned} - \int q_\sigma(\tilde{x}) \nabla_{\tilde{x}} \log q_\sigma(\tilde{x})^T s_\theta(\tilde{x}) d\tilde{x} &= - \int \nabla_{\tilde{x}} q_\sigma(\tilde{x})^T s_\theta(\tilde{x}) d\tilde{x} \\ &= - \int \nabla_{\tilde{x}} \left(\int p_{\text{data}}(x) q_\sigma(\tilde{x}|x) dx \right)^T s_\theta(\tilde{x}) d\tilde{x} \\ &= - \int \left(\int p_{\text{data}}(x) \nabla_{\tilde{x}} q_\sigma(\tilde{x}|x) dx \right)^T s_\theta(\tilde{x}) d\tilde{x} \\ &= - \int \int p_{\text{data}}(x) q_\sigma(\tilde{x}|x) \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x)^T s_\theta(\tilde{x}) dx d\tilde{x} \\ &= - \mathbb{E}_{x \sim p_{\text{data}}(x), \tilde{x} \sim q_\sigma(\tilde{x}|x)} \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x)^T s_\theta(\tilde{x}) \end{aligned}$$

- ▶ Plug it back we have

$$\begin{aligned} & \frac{1}{2} \mathbb{E}_{\tilde{x} \sim q_\sigma} \| \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}) - s_\theta(\tilde{x}) \|^2 \\ &= \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}(x), \tilde{x} \sim q_\sigma(\tilde{x}|x)} \| s_\theta(\tilde{x}) - \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x) \|^2 + \text{Const} \end{aligned}$$

- ▶ The noise score $\nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x)$ is easy to compute. For example, when use Gaussian noise $q_\sigma(\tilde{x}|x) = \mathcal{N}(\tilde{x}|x, \sigma^2 I)$, the score is

$$\nabla_{\tilde{x}} \log q_\sigma(\tilde{x}|x) = -\frac{\tilde{x} - x}{\sigma^2}$$

- ▶ Pros: efficient to optimize even for very high dimensional data, and useful for optimal denoising.
- ▶ Cons: cannot estimate the score of clean data (noise-free)

- ▶ Sample a minibatch of datapoints $\{x_1, \dots, x_n\} \sim p_{\text{data}}(x)$.
- ▶ Sample a minibatch of perturbed datapoints

$$\tilde{x}_i \sim q_\sigma(\tilde{x}_i|x_i), \quad i = 1, 2, \dots, n$$

- ▶ Estimate the denoising score matching loss with empirical means

$$\frac{1}{2n} \sum_{i=1}^n \|s_\theta(\tilde{x}) - \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}_i|x)\|^2$$

- ▶ Stochastic gradient descent
- ▶ Need to choose a very small σ ! However, the loss variance would also increase drastically as $\sigma \rightarrow 0$!

- ▶ Denoising score matching is suitable for optimal denoising
- ▶ Given $p(x)$, $q_\sigma(\tilde{x}|x) = \mathcal{N}(\tilde{x}|x, \sigma^2 I)$, we can define the posterior $p(x|\tilde{x})$ with Bayes' rule

$$p(x|\tilde{x}) = \frac{p(x)q_\sigma(\tilde{x}|x)}{q_\sigma(\tilde{x})}$$

where


$$q_\sigma(\tilde{x}) = \int p(x)q_\sigma(\tilde{x}|x)dx$$

- ▶ Tweedie's formula:

$$\begin{aligned}\mathbb{E}_{x \sim p(x|\tilde{x})}[x] &= \tilde{x} + \sigma^2 \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}) \\ &\approx \tilde{x} + \sigma^2 s_\theta(\tilde{x})\end{aligned}$$



- ▶ One dimensional problems should be easier.
- ▶ Consider projections onto random directions.
- ▶ Sliced score matching (Song et al 2019).



- ▶ Objective: Sliced Fisher Divergence

$$\frac{1}{2} \mathbb{E}_{v \sim p_v} \mathbb{E}_{x \sim p_{\text{data}}} (v^T \nabla_x \log p_{\text{data}}(x) - v^T s_\theta(x))^2$$

- ▶ Similarly, we can do integration by parts

$$\mathbb{E}_{v \sim p_v} \mathbb{E}_{x \sim p_{\text{data}}} \left(v^T \nabla_x s_\theta(x) v + \frac{1}{2} (v^T s_\theta(x))^2 \right)$$

- ▶ Computing Jacobian-vector products is scalable

$$v^T \nabla_x s_\theta(x) v = v^T \nabla_x (s_\theta(x)^T v)$$

This only requires one backpropagation!

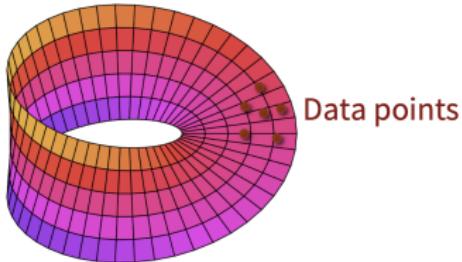


- ▶ Sample a minibatch of datapoints $\{x_1, \dots, x_n\} \sim p_{\text{data}}(x)$
- ▶ Sample a minibatch of projection directions $\{v_i \sim p_v\}_{i=1}^n$
- ▶ Estimate the sliced score matching loss with empirical means


$$\frac{1}{n} \sum_{i=1}^n \left(v_i^T \nabla_x s_\theta(x_i) v_i + \frac{1}{2} (v_i^T s_\theta(x_i))^2 \right)$$

- ▶ The perturbation distribution is typically Gaussian or Rademacher. When $\mathbb{E} vv^T = I$, this is equivalent to the **Hutchinson's trick**.
- ▶ Can use $\|s_\theta(x)\|^2$ instead of $(v^T s_\theta(x))^2$ to reduce variance.
- ▶ Can use more projections per datapoint to boost performance.

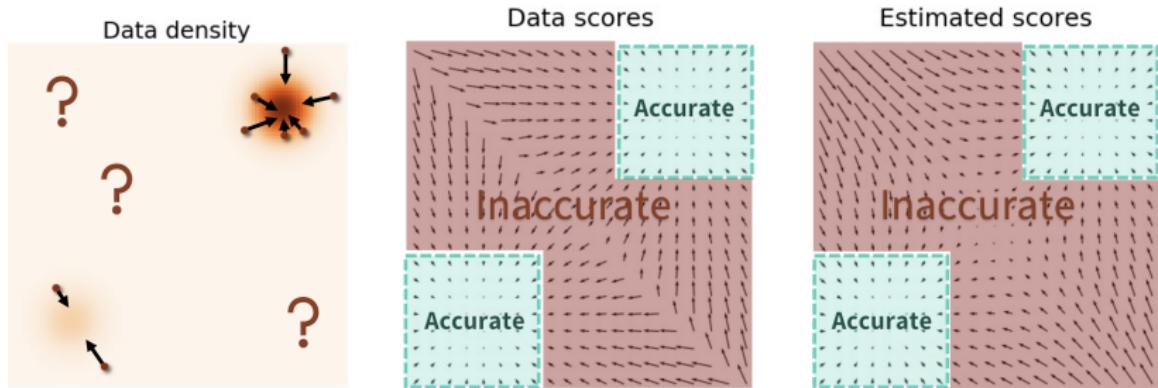
- ▶ Datapoints would lie on a lower dimensional manifold.



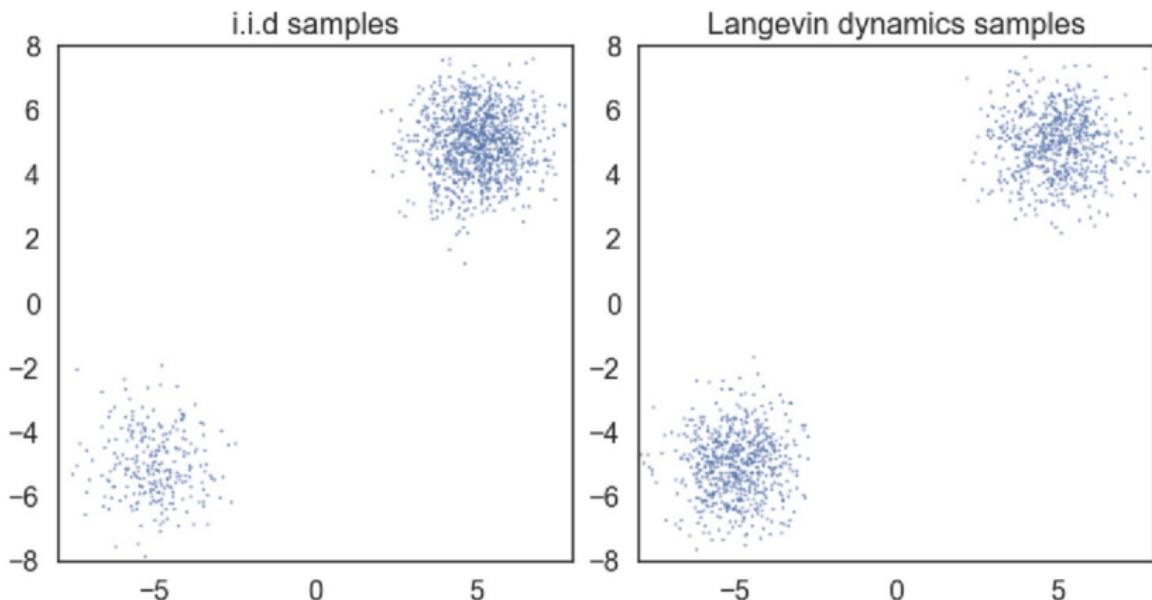
- ▶ Data score hence would be undefined.



$$\frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} \| \nabla_x \log p_{\text{data}}(x) - s_\theta(x) \|^2$$

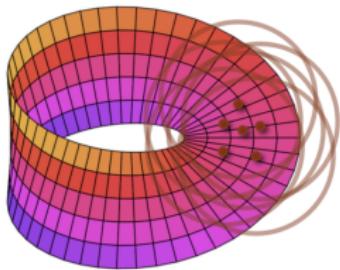


- ▶ Poor score estimation in low data density regions.
- ▶ Langevin MCMC will also have trouble exploring low density regions.

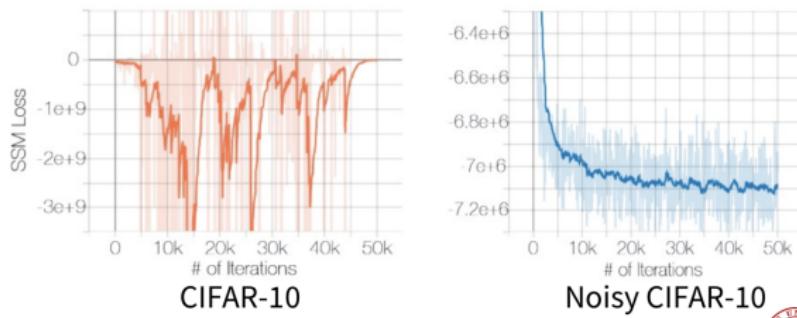


Adapted from Song et al 2019.

- ▶ The solution to all pitfalls: Gaussian perturbation!
- ▶ Inflate the flat manifold with noise.

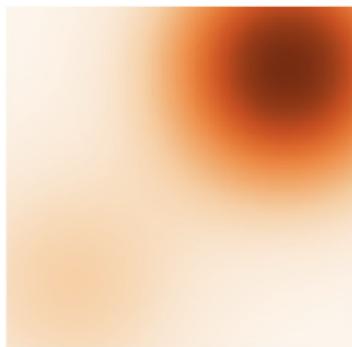


- ▶ Score matching on noise data

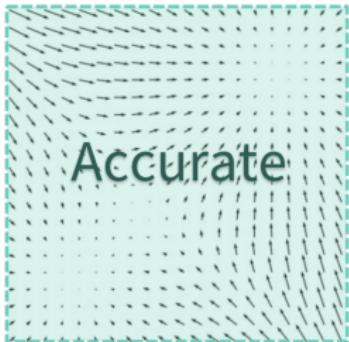


$$\frac{1}{2} \mathbb{E}_{\tilde{x} \sim q_\sigma} \| \nabla_{\tilde{x}} \log q_\sigma(\tilde{x}) - s_\theta(\tilde{x}) \|^2$$

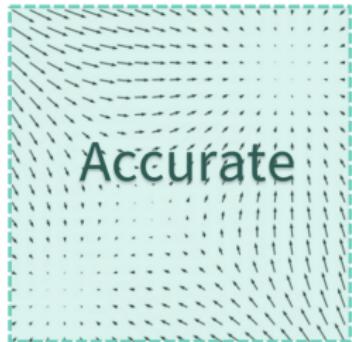
Perturbed density



Perturbed scores



Estimated scores




- ▶ Noisy score can provide useful directional information for Langevin MCMC.

- ▶ Multi-scale noise perturbations.



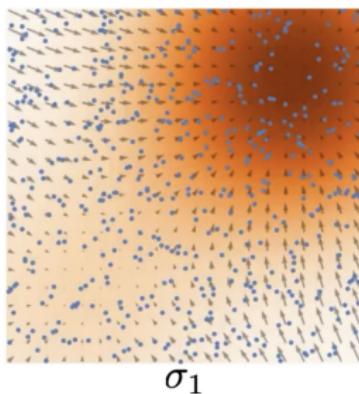
$$\sigma_1 > \sigma_2 > \dots > \sigma_{L-1} > \sigma_L$$



- ▶ Trading off data quality and estimator accuracy



- ▶ Sample using $\sigma_1, \dots, \sigma_L$ sequentially with Langevin dynamics.
- ▶ Anneal down the noise level.
- ▶ Samples used as initialization for the next level.




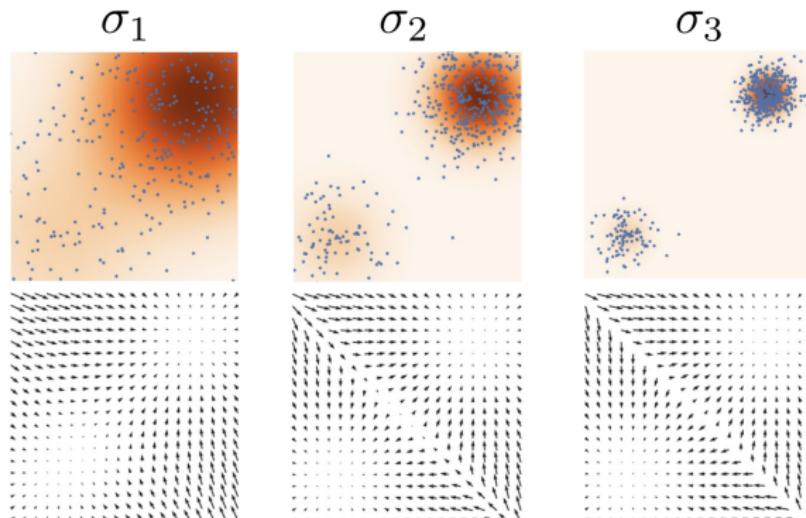
Algorithm 1 Annealed Langevin dynamics.

Require: $\{\sigma_i\}_{i=1}^L, \epsilon, T$.

- 1: Initialize $\tilde{\mathbf{x}}_0$
- 2: **for** $i \leftarrow 1$ to L **do**
- 3: $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$ ▷ α_i is the step size.
- 4: **for** $t \leftarrow 1$ to T **do**
- 5: Draw $\mathbf{z}_t \sim \mathcal{N}(0, I)$
- 6: $\tilde{\mathbf{x}}_t \leftarrow \tilde{\mathbf{x}}_{t-1} + \frac{\alpha_i}{2} \mathbf{s}_{\theta}(\tilde{\mathbf{x}}_{t-1}, \sigma_i) + \sqrt{\alpha_i} \mathbf{z}_t$
- 7: **end for**
- 8: $\tilde{\mathbf{x}}_0 \leftarrow \tilde{\mathbf{x}}_T$
- 9: **end for**
- return $\tilde{\mathbf{x}}_T$



- ▶ Learning score functions jointly with noise conditional score networks!



Noise Conditional
Score Network
(NCSN)

- ▶ As the goal is to estimate the score of perturbed data distributions, we can use denoising score matching for training.
- ▶ Assign different weights to combine denoising score matching losses for different noise levels.

$$\begin{aligned} & \frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) \mathbb{E}_{\tilde{x} \sim q_{\sigma_i}(\tilde{x})} \| \nabla_{\tilde{x}} \log q_{\sigma_i}(\tilde{x}) - s_{\theta}(\tilde{x}, \sigma_i) \|^2 \\ = & \frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) \mathbb{E}_{x \sim p_{\text{data}}, \tilde{x} \sim q_{\sigma_i}(\tilde{x}|x)} \| \nabla_x \log q_{\sigma_i}(\tilde{x}|x) - s_{\theta}(\tilde{x}, \sigma_i) \|^2 + \text{Const} \\ = & \frac{1}{L} \sum_{i=1}^L \lambda(\sigma_i) \mathbb{E}_{x \sim p_{\text{data}}, z \sim \mathcal{N}(0, I)} \| s_{\theta}(x + \sigma_i z, \sigma_i) + \frac{z}{\sigma_i} \|^2 + \text{Const}. \end{aligned}$$



- Adjacent noise scales should have sufficient overlap to ease transitioning across noise scales in annealed Langevin dynamics.
- For example, a geometric progression

$$\frac{\sigma_i}{\sigma_{i+1}} = \alpha > 1, \quad i = 1, \dots, L-1$$

- What about the weighting function λ ?
- Use $\lambda(\sigma) = \sigma^2$ to balance different score matching losses

$$\begin{aligned} & \frac{1}{L} \sum_{i=1}^L \sigma_i^2 \mathbb{E}_{x \sim p_{\text{data}}, z \sim \mathcal{N}(0, I)} \|s_\theta(x + \sigma_i z, \sigma_i) + \frac{z}{\sigma_i}\|^2 \\ &= \frac{1}{L} \sum_{i=1}^L \mathbb{E}_{x \sim p_{\text{data}}, z \sim \mathcal{N}(0, I)} \|\sigma_i s_\theta(x + \sigma_i z, \sigma_i) + z\|^2 \end{aligned}$$

- ▶ Sample a mini-batch of datapoints $\{x_1, \dots, x_n\} \sim p_{\text{data}}$.
- ▶ Sample a mini-batch of noise scale indices

$$\{i_1, \dots, i_n\} \sim \mathcal{U}\{1, 2, \dots, L\}$$

- ▶ Sample a mini-batch of Gaussian noise

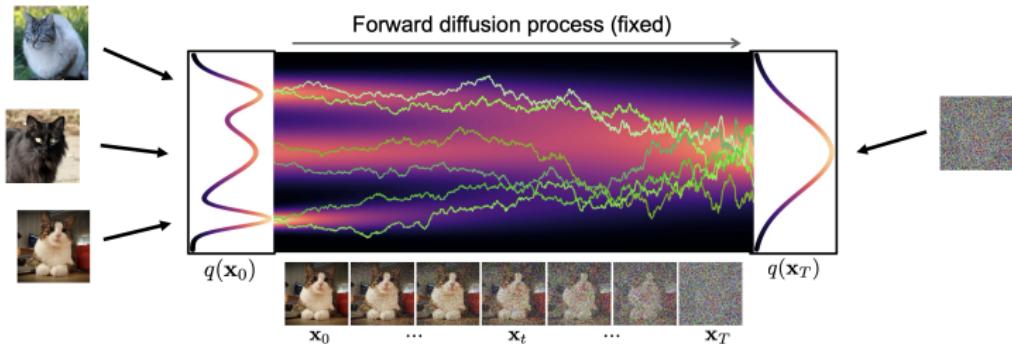
$$\{z_1, \dots, z_n\} \sim \mathcal{N}(0, I)$$

- ▶ Estimate the weighted mixture of score matching losses

$$\frac{1}{n} \sum_{k=1}^n \|\sigma_{i_k} s_\theta(x_k + \sigma_{i_k} z_k, \sigma_{i_k}) + z_k\|^2$$

- ▶ As efficient as training one single non-conditional score-based model.


Consider the case of infinitely many noise levels



Forward diffusion SDE: $dX_t = f(X_t, t)dt + g(t)dB_t.$
Examples:

- ▶ Variance Exploding: $f(X_t, t) = 0, \quad g(t) = \sqrt{\frac{d\sigma_t^2}{dt}}.$
- ▶ Variance Preserving: $f(X_t, t) = -X_t, \quad g(t) = \sqrt{2}.$





- ▶ Forward diffusion SDE:

$$d\bar{X}_t = f(\bar{X}_t, t)dt + g(t)dB_t, \quad \bar{X}_t \sim q_t.$$

- ▶ Reverse diffusion SDE: let $\bar{X}_t^\leftarrow := \bar{X}_{T-t}$, $0 \leq t \leq T$

$$d\bar{X}_t^\leftarrow = (g(T-t)^2 \nabla \log q_{T-t}(\bar{X}_t^\leftarrow) - f(\bar{X}_t^\leftarrow, T-t))dt + g(T-t)dB_t.$$

- ▶ Let q be the data distribution. Consider the OU forward process:

$$d\bar{X}_t = -\bar{X}_t dt + \sqrt{2} dB_t, \quad q_0 \sim q.$$

- ▶ The condition distribution is

$$\bar{X}_t | \bar{X}_0 \sim \mathcal{N}(e^{-t}\bar{X}_0, (1 - e^{-2t})I_d).$$

- ▶ The corresponding reverse process is

$$d\bar{X}_t^\leftarrow = (\bar{X}_t^\leftarrow + 2\nabla \log q_{T-t}(\bar{X}_t^\leftarrow))dt + \sqrt{2} dB_t.$$

where q_t is the law of the forward process.

- ▶ Denoising score matching:

$$\min_s \mathbb{E}_{\bar{x}_0 \sim q, \bar{x}_t \sim q(\bar{x}_t | \bar{x}_0)} \|s_t(\bar{x}_t) - \nabla_{\bar{x}_t} \log q(\bar{x}_t | \bar{x}_0)\|^2.$$

- ▶ Reverse SDE with estimated score

$$d\bar{X}_t^\leftarrow = (\bar{X}_t^\leftarrow + 2s_{T-t}(\bar{X}_t^\leftarrow))dt + \sqrt{2}dB_t.$$

- ▶ Let $h > 0$ be the step size. Assume that we have score estimates s_{kh} for each time $k = 0, 1, \dots, N$, where $T = Nh$.
- ▶ Discretize the reverse SDE using an exponential integrator

$$d\bar{X}_t^\leftarrow = (\bar{X}_t^\leftarrow + 2s_{T-kh}(\bar{X}_{kh}^\leftarrow))dt + \sqrt{2}dB_t, \quad t \in [kh, (k+1)h]$$

- ▶ How well can the data distribution be approximated if the score estimation is accurate enough?

Assumptions:

- ▶ A1: $\forall t \geq 0$, the score function $\nabla \log q_t$ L -Lipschitz.
- ▶ A2: For some $\eta > 0$, $\mathbb{E}_q \|\cdot\|^{2+\eta}$ is finite, and

$$m_2^2 := \mathbb{E}_q \|\cdot\|^2.$$

- ▶ A3: For all $k = 1, N$, $\mathbb{E}_{q_{kh}} \|s_{kh} - \nabla \log q_{kh}\|^2 \leq \epsilon^2$.

Theorem (Chen et al., 2023)

Suppose A1-3 hold. Let p_T be the output of the discretized reverse SDE at time T with $\bar{X}_0^\leftarrow \sim \gamma^d$, and suppose $h \lesssim 1/L$, where $L \geq 1$. Then it holds that

$$\text{TV}(p_T, q) \lesssim \sqrt{\text{KL}(q\|\gamma^d)} \exp(-T) + (L\sqrt{dh} + Lm_2h)\sqrt{T} + \epsilon\sqrt{T}$$

- ▶ Let Q_T^\leftarrow be the path measure of the exact reverse process

$$d\bar{X}_t^\leftarrow = (\bar{X}_t^\leftarrow + 2\nabla \log q_{T-t}(\bar{X}_t^\leftarrow))dt + \sqrt{2}dB_t.$$

- ▶ Let $P_T^{q_T}$ be the path measure of the approximated reverse process

$$d\bar{X}_t^\leftarrow = (\bar{X}_t^\leftarrow + 2s_{T-kh}(\bar{X}_{kh}^\leftarrow))dt + \sqrt{2}dB_t, \quad t \in [kh, (k+1)h]$$

- ▶ Girsanov's theorem: a more general case

$$\text{KL}(Q_T^\leftarrow \| P_T^{q_T}) \leq \sum_{k=0}^{N-1} \mathbb{E}_{Q_T^\leftarrow} \int_{kh}^{(k+1)h} \| s_{T-kh}(X_{kh}) - \nabla \log q_{T-t}(X_t) \|^2 dt.$$



- Bounding the discretization error. $\forall t \in [kh, (k+1)h]$

$$\begin{aligned} & \mathbb{E}_{Q_T^\leftarrow} \|s_{T-kh}(X_{kh}) - \nabla \log q_{T-t}(X_t)\|^2 \\ & \lesssim \mathbb{E}_{Q_T^\leftarrow} (\|s_{T-kh}(X_{kh}) - \nabla \log q_{T-kh}(X_{kh})\|^2 \\ & \quad + \|\nabla \log q_{T-kh}(X_{kh}) - \nabla \log q_{T-t}(X_{kh})\|^2 \\ & \quad + \|\nabla \log q_{T-t}(X_{kh}) - \nabla \log q_{T-t}(X_t)\|^2) \\ & \lesssim \epsilon^2 + \mathbb{E}_{Q_T^\leftarrow} \left\| \nabla \log \frac{q_{T-kh}}{q_{T-t}}(X_{kh}) \right\|^2 + L^2 \mathbb{E}_{Q_T^\leftarrow} \|X_{kh} - X_t\|^2. \end{aligned}$$



- ▶ Bounding the change of score along the forward process

$$\left\| \nabla \log \frac{q_{T-kh}}{q_{T-t}}(X_{kh}) \right\|^2 \lesssim L^2 dh + L^2 h^2 \|X_{kh}\|^2 + L^2 h^2 \|\nabla \log q_{T-t}(X_{kh})\|^2.$$

- ▶ For the last term

$$\begin{aligned} \|\nabla \log q_{T-t}(X_{kh})\|^2 &\lesssim \|\nabla \log q_{T-t}(X_t)\|^2 + \\ &\quad \|\nabla \log q_{T-t}(X_{kh}) - \nabla \log q_{T-t}(X_t)\|^2 \\ &\lesssim \|\nabla \log q_{T-t}(X_t)\|^2 + L^2 \|X_{kh} - X_t\|^2. \end{aligned}$$

- ▶ put these together

$$\begin{aligned} & \mathbb{E}_{Q_T^\leftarrow} \|s_{T-kh}(X_{kh}) - \nabla \log q_{T-t}(X_t)\|^2 \\ & \lesssim \epsilon^2 + L^2 dh + L^2 h^2 \mathbb{E}_{Q_T^\leftarrow} \|X_{kh}\|^2 \\ & \quad + L^2 h^2 \mathbb{E}_{Q_T^\leftarrow} \|\nabla \log q_{T-t}(X_t)\|^2 + L^2 \mathbb{E}_{Q_T^\leftarrow} \|X_{kh} - X_t\|^2. \end{aligned}$$

- ▶ apply moment bounds for the forward process

$$\begin{aligned} & \mathbb{E}_{Q_T^\leftarrow} \|s_{T-kh}(X_{kh}) - \nabla \log q_{T-t}(X_t)\|^2 \\ & \lesssim \epsilon^2 + L^2 dh + L^2 h^2(d + m_2^2) + L^3 h^2 d + L^2(m^2 h^2 + dh) \\ & \lesssim \epsilon^2 + L^2 dh + L^2 m_2^2 h^2. \end{aligned}$$



- ▶ According to Girsanov's theorem

$$\text{KL}(Q_T^\leftarrow \| P_T^{q_T}) \lesssim (\epsilon^2 + L^2 dh + L^2 m_2^2 h^2)T.$$

- ▶ By data processing inequality

$$\begin{aligned}\text{TV}(p_T, q) &\leq \text{TV}(p_T^{\gamma^d}, p_T^{q_T}) + \text{TV}(p_T^{q_T}, Q_T^\leftarrow) \\ &\leq \text{TV}(q_T, \gamma^d) + \text{TV}(p_T^{q_T}, Q_T^\leftarrow).\end{aligned}$$

- ▶ Using the convergence of the OU process in KL divergence and Pinsker inequality, we have

$$\text{TV}(p_T, q) \lesssim \sqrt{\text{KL}(q\|\gamma^d)} \exp(-T) + (L\sqrt{dh} + Lm_2h)\sqrt{T} + \epsilon\sqrt{T}$$



- ▶ Salakhutdinov, R. and Hinton, G. Deep boltzmann machines. In Artificial intelligence and statistics, 2009.
- ▶ Aapo Hyvarinen. Estimation of non-normalized statistical models by score matching. Journal of Machine Learning Research, 6(Apr):695–709, 2005.
- ▶ Pascal Vincent. A connection between score matching and denoising autoencoders. Neural computation, 23(7):1661–1674, 2011.



- ▶ Yang Song, Sahaj Garg, Jiaxin Shi, and Stefano Ermon. Sliced score matching: A scalable approach to density and score estimation. In Proceedings of the Thirty-Fifth Conference on Uncertainty in Artificial Intelligence, UAI 2019.
- ▶ Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution. In Advances in Neural Information Processing Systems, pp. 11895–11907, 2019.
- ▶ Sitan Chen, Sinho Chewi, Jerry Li, Yuanzhi Li, Adil Salim, and Anru Zhang. Sampling is as easy as learning the score: theory for diffusion models with minimal data assumptions. In The Eleventh International Conference on Learning Representations, 2023

